

SentinelAI

Predict • Prevent • Protect

AI-Powered Digital Public Safety Intelligence Platform

Defeating Counterfeiting, Fraud & Digital Arrest Scams

ET AI Hackathon 2026

Theme: Smart Cities / Public Safety / Digital Trust / Geospatial Law Enforcement

Solution Document | Version 1.0

1. Executive Summary

India recorded 1.14 million cybercrime complaints in 2023 — a 60% year-on-year surge. Digital arrest scams alone drained Rs 1,776 crore from citizens in the first nine months of 2024, while the RBI's Annual Report 2025 flagged record seizures of high-quality counterfeit Rs 500 notes. These threats are no longer opportunistic; they are industrialised operations run from cross-border fraud compounds using spoofed telephony, AI-generated voices, and cloned government portals.

SentinelAI is an AI-powered Digital Public Safety Intelligence platform that unifies computer vision, graph analytics, NLP, speech AI, and geospatial intelligence into a single command surface for law enforcement, financial institutions, and citizens. It shifts the model from reactive complaint-taking to predictive threat neutralisation — detecting scams at the point of contact, mapping fraud networks before mass victimisation, and identifying counterfeit currency at the teller window.

This document describes the problem, target users, solution architecture, AI models, data flows, security posture, deployment strategy, business impact, and roadmap for SentinelAI as submitted to the ET AI Hackathon 2026.

Key Outcomes

- Detect digital arrest scams in-session with >92% precision and neutralise fund transfer before it occurs.
- Identify counterfeit currency across denominations with >95% accuracy on mobile-grade cameras.
- Reduce time-to-lead for organised fraud networks from weeks to minutes via graph clustering.
- Provide a citizen-facing multilingual shield with a false-positive rate below 2%.
- Deliver court-admissible intelligence packages with full audit trails and chain of custody.

2. Problem Context

2.1 The Scale

The Indian cybercrime landscape has crossed the threshold from nuisance to systemic public safety crisis. Three trends define the current moment:

- **Volume:** 1.14 million cybercrime complaints in 2023, up 60% YoY, with 2024 tracking materially higher.
- **Loss:** Rs 1,776 crore lost to digital arrest scams in just nine months of 2024 (MHA data).
- **Sophistication:** Record FICN seizures of Rs 500 notes of sufficient print quality to defeat manual detection at bank counters.

2.2 Why Existing Tools Fail

- Law enforcement receives evidence after victimisation — the money has already moved through mule networks.
- Bank counter staff rely on tactile checks that fail against modern intaglio-quality counterfeits.
- Citizens have no in-the-moment verification when a caller claims to be from CBI, ED, or Customs.
- Fraud intelligence is siloed across telecom, banking, and police jurisdictions — coordination happens after harm.

2.3 What Is Actually Needed

The gap is not evidentiary — it is intelligence before mass victimisation, and reliable point-of-contact detection. This requires convergence of four traditionally separate signal domains: financial transaction telemetry, communications metadata, physical security (currency), and geospatial law enforcement operations. That convergence is precisely the class of problem where multi-modal AI is transformative.

3. Solution Overview — SentinelAI

SentinelAI is a modular intelligence platform with five tightly integrated pillars, each addressing one dimension of the problem, all sharing a common event bus, identity graph, and audit ledger.

3.1 The Five Pillars

Pillar	Primary User	Core Capability
Digital Arrest Detection	Citizens, Telecom, MHA	Real-time classifier on call flow, spoof signatures, and script templates.
Counterfeit Currency Agent	Bank tellers, Police, Retail	On-device CV model verifying watermark, thread, microprint, latent image.
Fraud Network Graph	LEA investigators, Banks	Graph AI clustering transactions, devices, UPI handles, and mule accounts.
Geospatial Command Centre	State police, MHA	Live map of complaints, seizures, and hotspots for patrol optimisation.
Citizen Fraud Shield	Public (12 languages)	WhatsApp / IVR / app conversational agent for instant risk verdict.

3.2 Design Principles

- **Predictive over reactive:** Every module produces leads before financial transfer or mass victimisation.
- **Explainable by default:** Every verdict ships with a signal breakdown suitable for a court affidavit.
- **Low false-positive discipline:** Citizen-facing thresholds tuned for FPR < 2%; investigator-facing tuned for recall.
- **Federated by design:** Bank, telecom, and police data stay in their jurisdictions — only hashed features cross the wire.
- **Auditability:** Every model call, evidence artefact, and analyst action is logged to an append-only ledger.

4. Module Deep Dive

4.1 Digital Arrest Scam Detection

Threat Pattern

A caller impersonating CBI, ED, Customs, or Police claims the victim's Aadhaar or a parcel is linked to money laundering or narcotics. The victim is held on a video call, forbidden from disconnecting, and coerced into transferring funds to an 'RBI verification account'. Sessions can last hours to days.

AI Approach

- Speech-to-text pipeline (Whisper-family, IndicASR) transcribes call audio in real time.
- Transformer classifier trained on labelled scam transcripts detects impersonation, urgency, isolation, and payment coercion tokens.
- Voice biometrics module flags AI-generated voices via spectral and prosodic features (deepfake detection).
- Telecom metadata analyser correlates spoofed CLI, VoIP entry points, and repeat-caller device fingerprints.

Explainable Output

Each verdict ships with a signal panel: pattern match (CBI / ED / Customs / Courier), red-flag keywords, spoof indicators, and a confidence percentage. The verdict triggers three actions: warn the victim on-device, alert the telecom provider, and generate an MHA advisory.

Target Metrics

- **Precision:** > 92% on live-session detection
- **Recall:** > 88% within first 90 seconds of the call
- **Time-to-alert:** < 5 seconds from detection to victim notification

4.2 Counterfeit Currency Identification Agent

Threat Pattern

High-quality Rs 500 and Rs 2000 fakes clear tactile inspection at bank counters. Manual UV lamps miss subtle intaglio and microprint defects. Field officers lack any tool at the point of seizure.

AI Approach

- Convolutional backbone (EfficientNet / MobileViT) trained on multi-denomination note images at multiple lighting conditions.

- Feature-specific heads validate security thread, watermark (Gandhi portrait), microprint 'RBI', latent image, intaglio raised print, and serial-number pattern.
- Serial-number OCR cross-checks against RBI's issued-range database for pattern anomalies.
- On-device inference via TensorFlow Lite / ONNX Runtime Mobile — under 200ms on mid-range Android.

Deployment Surfaces

- **Mobile app:** Field officers, retail cashiers, citizens verifying received notes.
- **Bank counting machine SDK:** Real-time flagging inline with existing note counters.
- **POS terminal:** Retail counterfeit interception at the moment of tender.

Target Metrics

- **Accuracy:** > 95% across Rs 100, Rs 200, Rs 500, Rs 2000
- **Inference latency:** < 200ms on Snapdragon 6-gen mobile
- **False accept rate:** < 1.5% on high-quality fakes

4.3 Fraud Network Graph Intelligence

Threat Pattern

A single scam campaign spans hundreds of victims, dozens of mule accounts, shared devices, and a handful of controllers. Investigators today see one complaint at a time and miss the network.

AI Approach

- Heterogeneous graph of victims, phones, IMEI/IMSI, UPI handles, bank accounts, IPs, and geolocations.
- Graph Neural Network (GraphSAGE / RGCN) computes node embeddings that cluster co-orchestrated fraud.
- Community detection (Louvain, Leiden) surfaces mule rings and controller nodes.
- Temporal edge weighting prioritises fresh money flow within investigation windows.

Investigator UX

- Interactive graph with money-flow, communication, and relationship edges.
- Cluster cards summarising victims, common devices, mule-bank convergence, and estimated exposure.
- Export to a court-admissible PDF intelligence package with hashes, timestamps, and analyst attestation.

4.4 Geospatial Crime Pattern Intelligence

A live command centre map overlays complaint density, counterfeit seizure points, fraud call origination cells, and predicted hotspots. State-level dashboards support resource deployment and inter-district intelligence sharing in near real time.

- H3 hexagonal binning for uniform density comparison across urban and rural cells.
- Spatio-temporal forecasting (ST-GNN / Prophet) predicts hotspot migration 24-72 hours ahead.
- Patrol optimiser suggests beat allocation to minimise expected loss under a fixed patrol budget.

4.5 Citizen Fraud Shield


A conversational AI accessible over WhatsApp, IVR, and the SentinelAI mobile app that lets any citizen paste an SMS, forward a call recording, share a QR image, or describe a situation. The bot returns an instant verdict, an explanation in plain language, and a guided flow to the NCRB 1930 portal.

- Twelve regional languages supported via Indic-LLM adapters.
- Zero-PII inference — messages are hashed on-device for aggregate learning.
- Escalation to human counsellors for high-distress victims of ongoing scams.

5. System Architecture

5.1 High-Level Diagram

The platform is organised into five layers — Ingest, Feature, Model, Intelligence, and Delivery — connected by an event bus and governed by a shared identity graph and audit ledger.



5.2 Technology Stack

Layer	Technology	Rationale
Frontend	React 19 + TanStack Start, Tailwind v4, Framer Motion	SSR-capable, type-safe routing, premium UX.
API / BFF	TanStack server functions on Cloudflare Workers	Edge latency, zero-cold-start for citizen surface.
Data	Postgres (OLTP), ClickHouse (analytics), Neo4j (graph)	Right tool per workload; polyglot persistence.
ML Serving	Triton / TorchServe on GPU pool; ONNX for on-device	Server + edge coverage with one model artefact.
Streaming	Kafka + Flink for real-time enrichment	Sub-second alerting on live scam sessions.
Identity	OIDC + role-based access (Citizen / Officer / Admin)	Least-privilege, MFA for LEA users.
Storage	S3-compatible object store with WORM ledger	Court-admissible evidence integrity.

5.3 Data Flow — Digital Arrest Detection

- Telecom mirror stream ingests call metadata into Kafka.
- Citizen app (or telecom edge) streams transcribed segments to the Scam NLP model.
- Model returns verdict + signal panel within 5s; verdict enters the Alerts queue.
- Alerts fan out to: citizen device (banner + interrupt), telecom SOC, MHA dashboard.
- Every artefact hashed and written to the audit ledger with signed timestamps.

6. Security, Privacy & Compliance

6.1 Threat Model

- Adversary 1: fraud operators attempting to poison training data or bypass detection.
- Adversary 2: insider misuse of citizen data by authorised officers.
- Adversary 3: nation-state exfiltration of the identity graph.

6.2 Controls

- **Data minimisation:** Only hashed identifiers cross agency boundaries; raw PII stays in source-of-truth.
- **Encryption:** TLS 1.3 in transit; AES-256 at rest; KMS-backed per-tenant keys.
- **Access control:** OIDC + MFA; RBAC with attribute-based scoping per district and case.
- **Audit ledger:** Append-only, hash-chained, exportable as a signed evidentiary bundle.
- **Model integrity:** Signed model artefacts, canary evaluation, drift detection, adversarial red-team pipeline.
- **Compliance:** DPDP Act 2023 alignment, CERT-In directives, MHA data-sharing protocols.

6.3 Legal Admissibility

Every intelligence package generated for LEA use includes: chain-of-custody hashes, model version, feature snapshot, analyst attestation, and reproducibility metadata. This makes the package suitable for submission under the Indian Evidence Act as electronic evidence.

7. User Experience

SentinelAI is designed to look and feel like a premium platform product — the standard set by Microsoft, Google, Palantir, and Cisco — while remaining approachable for a first-time citizen user.

7.1 Citizen Surface

- Landing page in the citizen's preferred language with a single primary action: 'Check a message, call, or note'.
- Guided reporting to the NCRB 1930 portal with pre-filled fields and case ID.
- Real-time SMS / call scanning with a plain-language verdict card.

7.2 Officer Surface

- Dark-mode command centre with map, alert feed, and case queue.
- Graph explorer with money-flow animations and cluster insights.
- One-click intelligence package export.

7.3 Bank / Retail Surface

- Currency scanner with instant verdict and feature-level breakdown.
- Teller override with mandatory reason for audit trail.

8. Business Impact

8.1 Quantified Impact (Year 1 target)

Outcome	Baseline	SentinelAI Target
Digital arrest loss per victim	Rs 22 lakh (avg)	Rs 0 (intercepted pre-transfer) in >70% of flagged sessions
Counterfeit interception at bank	~15% (manual)	>90% (AI-assisted)
Time from first complaint to network map	3-6 weeks	< 30 minutes
Citizen tool false positive rate	N/A	< 2%
Court-admissible packages / month	Manual, weeks each	Automated, minutes each

8.2 Stakeholder Value

- **Citizens:** In-the-moment protection in their own language; guided recovery when victimised.
- **Law enforcement:** Predictive leads, network visibility across jurisdictions, admissible evidence.
- **Banks & telecom:** Reduced fraud exposure, regulatory alignment, safer customer journeys.
- **MHA / RBI:** National situational awareness, faster policy feedback loops.

9. Scalability & Deployment

- Multi-tenant SaaS for LEA states with per-tenant data isolation and KMS keys.
- Edge inference on citizen devices and bank machines — server load scales with alerts, not users.
- Kafka + Flink pipelines handle millions of CDR and transaction events per minute.
- Graph shard strategy scales Neo4j horizontally by district for LEA queries.
- Cloudflare Workers front-end achieves sub-100ms TTFB globally for the citizen surface.

10. Roadmap

Phase	Timeline	Milestones
MVP (Hackathon)	Complete	5 modules live, mock intelligence engine, premium UX.
Pilot	Q1 2026	One state LEA + one PSU bank; real telco and transaction feeds.
Scale	Q2-Q3 2026	Five states, 3 banks, WhatsApp Business API for citizen bot.
National	Q4 2026	MHA / NCRB integration; RBI counterfeit reporting standard.

11. Alignment with Evaluation Focus

- **Counterfeit detection accuracy:** Feature-level heads deliver >95% across denominations; UV / IR extension planned.
- **Digital arrest precision & recall:** Multi-modal (text + voice + metadata) fusion targets >92% precision, >88% recall.
- **Fraud network lead time:** Streaming graph updates cut lead time from weeks to minutes.
- **Citizen-facing FPR:** Conservative thresholds keep FPR under 2%; every verdict is explainable.
- **Auditability:** Append-only ledger + signed intelligence packages ensure legal admissibility.

Judging Criteria Coverage

Criterion	Weight	SentinelAI Strength
Innovation	25%	Multi-modal fusion of CV, NLP, Voice, Graph, Geo in one platform.
Business Impact	25%	Direct reduction in Rs 1,776 cr+ annual loss; measurable per-module KPIs.
Technical Excellence	20%	Edge + cloud inference, GNN, streaming, court-grade audit trail.
Scalability	15%	Federated architecture; edge inference decouples users from server cost.
User Experience	15%	Premium command centre; 12-language citizen surface; sub-second verdicts.

12. Deliverables

- **Working Prototype:** Live SentinelAI web application covering all five pillars with functional mock intelligence.
- **Architecture Diagram:** Included in Section 5; expanded C4-style diagrams in the appendix.
- **Presentation Deck:** 12-slide deck aligned to judging criteria.
- **Demo Video:** 3-minute walkthrough covering citizen, officer, and bank journeys.

13. Team & Roles

SentinelAI is delivered by a cross-functional team acting as senior software architects, UI/UX designers, frontend engineers, backend engineers, AI/ML engineers, and cybersecurity experts — with domain input modelled on published MHA, RBI, and CERT-In advisories.

14. Conclusion

Digital public safety in India is at an inflection point. The volume, sophistication, and industrialisation of digital fraud and counterfeit currency have outpaced tools designed for a slower, more local threat landscape. SentinelAI closes that gap by fusing computer vision, NLP, speech AI, graph analytics, and geospatial intelligence into a single, auditable, citizen-friendly platform — moving India's public safety posture from reactive investigation to predictive protection.

Predict. Prevent. Protect.